

Syllabus etc.

- Respond to my **Slack** invitation. Examine the board: <https://stat-s782fa2017.slack.com/>.
- If you don't have a **github** account, set one up: <https://education.github.com/pack>.
- Consult the course website: <https://github.com/stats-782fa2017>.
- Start thinking about a paper you want to present (due on Thursday, 24 August) and a group you want for your project.

1.1 Random variables

A *random variable* is a map X from a probability space Ω to \mathbb{R} . We write

$$P(X \in A) = P(\{\omega \in \Omega : X(\omega) \in A\}) \quad (1.1)$$

and we write $X \sim P$ to mean that X has distribution P .

The *cumulative distribution function* (cdf) of X is

$$F_X(x) = F(x) = P(X \leq x). \quad (1.2)$$

If X is discrete, its probability mass function (pmf) is

$$p_X(x) = p(x) = P(X = x). \quad (1.3)$$

If X is continuous, then its probability density function (pdf) satisfies

$$P(X \in A) = \int_A p_X(x) dx = \int_A p(x) dx \quad (1.4)$$

and $p_X(x) = p(x) = F'(x)$. The following are all equivalent:

$$X \sim P, \quad X \sim F, \quad X \sim p, \quad \mathcal{L}(X) = P. \quad (1.5)$$

Suppose that $X \sim P$ and $Y \sim Q$. We say that X and Y have the same distribution if

$$P(X \in A) = Q(Y \in A) \quad (1.6)$$

for all A . In other words, $P = Q$. In that case we say that X and Y are equal in distribution and we write $X \stackrel{d}{=} Y$ or $\mathcal{L}(X) = \mathcal{L}(Y)$. It can be shown that $X \stackrel{d}{=} Y$ if and only if $F_X(t) = F_Y(t)$ for all t .

1.2 Expected values

The *mean* or expected value of $g(X)$ is

$$\mathbb{E}[g(X)] = \int g(x)dF_X(x) = \int g(x)dP(x) = \begin{cases} \int_{-\infty}^{\infty} g(x)p(x)dx & \text{if } X \text{ is continuous} \\ \sum_j g(x_j)p(x_j) & \text{if } X \text{ is discrete.} \end{cases} \quad (1.7)$$

Recall the following useful properties of expectation:

1. $\mathbb{E}\left[\sum_{j=1}^k c_j g_j(X)\right] = \sum_{j=1}^k c_j \mathbb{E}[g_j(X)]$.
2. If X_1, \dots, X_n are independent then $\mathbb{E}[\prod_{i=1}^n X_i] = \prod_i \mathbb{E}[X_i]$.
3. We often write $\mu = \mathbb{E}[X]$.
4. $\sigma^2 = \text{Var}[X] = \mathbb{E}[(X - \mu)^2]$ is the Variance.
5. $\text{Var}[X] = \mathbb{E}[X^2] - \mu^2$.
6. If X_1, \dots, X_n are independent then $\text{Var}[\sum_{i=1}^n a_i X_i] = \sum_i a_i^2 \text{Var}[X_i]$
7. The covariance is $\text{Cov}[X, Y] = \mathbb{E}[(X - \mu_X)(Y - \mu_Y)] = \mathbb{E}[XY] - \mu_X \mu_Y$ and the correlation is $\rho(X, Y) = \text{Cov}[X, Y] / \sigma_X \sigma_Y$. Recall that $-1 \leq \rho(X, Y) \leq 1$.

The conditional expectation of Y given X is the random variable $\mathbb{E}[Y|X]$ whose value, when $X = x$ is $\mathbb{E}[Y|X = x] = \int yp(y|x)dy$ where $p(y|x) = p(x, y)/p(x)$. The *Law of Total Expectation* or *Law of Iterated Expectation*:

$$\mathbb{E}[Y] = \mathbb{E}[\mathbb{E}[Y|X]] = \int \mathbb{E}[Y|X = x] p_X(x) dx. \quad (1.8)$$

1.3 Important distributions

Normal $X \sim N(\mu, \sigma^2)$ if

$$p(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{1}{2\sigma^2}(x - \mu)^2\right\} \quad (1.9)$$

Multivariate normal For $X \in \mathbb{R}^d$, $X \sim N_d(\mu, \Sigma)$ if

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp\left\{-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)\right\} \quad (1.10)$$

Bernoulli $X \sim \text{Bernoulli}(\theta)$ if $P(X = 1) = \theta$ and $P(X = 0) = 1 - \theta$. Thus the pdf is

$$p(x; \theta) = \theta^x (1 - \theta)^{1-x} I_{\{0,1\}}(x). \quad (1.11)$$

Binomial $X \sim \text{Binomial}(\theta)$ if

$$p(x; \theta) = P(X = x) = \binom{n}{x} \theta^x (1 - \theta)^{n-x} I_{\{0, \dots, n\}}(x). \quad (1.12)$$

Uniform $X \sim U(a, b)$ if $p(x) = I_{[a,b]}(x)/(b - a)$.

1.4 Convergence

Much of statistical machine learning is concerned with showing asymptotic, or better yet, finite sample properties of computational methods. We want to know whether our methods will perform well, and if so how well will they perform relative to other methods. For that, we need to discuss some useful results from probability theory. First, we'll review the convergence of random variables.

Let X_1, X_2, \dots be a sequence of random variables, and let X be another random variable with distribution P . Let F_n be the cdf of X_n and let F be the cdf of X .

1. X_n converges *almost surely* to X , $X_n \xrightarrow{as} X$, if for every $\epsilon > 0$,

$$P\left(\lim_{n \rightarrow \infty} |X_n - X| < \epsilon\right) = 1.^1 \quad (1.13)$$

2. X_n converges *in probability* to X , $X_n \xrightarrow{P} X$, if for every $\epsilon > 0$,

$$\lim_{n \rightarrow \infty} P(|X_n - X| < \epsilon) = 1. \quad (1.14)$$

3. X_n converges *in L_p* to X , $X_n \xrightarrow{L_p} X$, if

$$\lim_{n \rightarrow \infty} \int |X_n - X|^p dP = \lim_{n \rightarrow \infty} \mathbb{E}[|X_n - X|^p] = 0. \quad (1.15)$$

4. X_n converges *in distribution* to X , $X_n \rightsquigarrow X$, if

$$\lim_{n \rightarrow \infty} F_n(t) = F(t) \quad (1.16)$$

for all t for which F is continuous.

Example 1.1. *This example shows that convergence in probability does not imply almost sure convergence. Let $S = [0, 1]$. Let P be uniform on $[0, 1]$. We draw SP . Let $X(s) = s$ and let*

$$\begin{array}{lll} X_1 = s + I_{[0,1]}(s) & X_2 = s + I_{[0,1/2]}(s) & X_3 = s + I_{[1/2,1]}(s) \\ X_4 = s + I_{[0,1/3]}(s) & X_5 = s + I_{[1/3,2/3]}(s) & X_6 = s + I_{[2/3,1]}(s) \end{array}$$

etc. Then $X_n \xrightarrow{P}$ since $P(|X_n - X| > \epsilon)$ is equal to the probability of an interval of s values whose length is going to zero. However, for every s , $X_n(s)$ alternates between the values s and $s + 1$ infinitely often, so this convergence does not occur almost surely.

Theorem 1.2. *The following relationships hold:*

- (a) $X_n \xrightarrow{L_p} X$ implies that $X_n \xrightarrow{P} X$.
- (b) $X_n \xrightarrow{P} X$ implies that $X_n \rightsquigarrow X$.
- (c) If $X_n \rightsquigarrow X$ and if $P(X = c) = 1$ for some real number c , then $X_n \xrightarrow{P} X$.
- (d) $X_n \xrightarrow{as} X$ implies that $X_n \xrightarrow{P} X$.

Theorem 1.3. *Let X_n, X, Y_n, Y be random variables. Let g be a continuous function. Let c be a constant.*

- (a) *If $X_n \xrightarrow{P} X$ and $Y_n \xrightarrow{P} Y$, then $X_n + Y_n \xrightarrow{P} X + Y$.*

¹The absolute value here can be replaced by any appropriate distance.

- (b) If $X_n \xrightarrow{L_p} X$ and $Y_n \xrightarrow{L_p} Y$, then $X_n + Y_n \xrightarrow{L_p} X + Y$.
- (c) If $X_n \rightsquigarrow X$ and $Y_n \rightsquigarrow c$, then $X_n + Y_n \rightsquigarrow X + c$.
- (d) If $X_n \xrightarrow{P} X$ and $Y_n \xrightarrow{P} Y$, then $X_n Y_n \xrightarrow{P} XY$.
- (e) If $X_n \rightsquigarrow X$ and $Y_n \rightsquigarrow c$, then $X_n Y_n \rightsquigarrow cX$.
- (f) If $X_n \xrightarrow{\text{converges somehow}} X$, then $g(X_n) \xrightarrow{\text{converges the same}} g(X)$.

- Parts (c) and (e) are known as Slutsky's theorem.
- Part (f) is known as The continuous mapping theorem.

1.5 LLNs and CLT

Theorem 1.4 (Weak law of large numbers). *Let X_1, X_2, \dots be independent random variables, each with finite mean and variance. Define $S_n = X_1 + \dots + X_n$. Then $\frac{1}{n}(S_n - \mathbb{E}[S_n]) \xrightarrow{P} 0$.*

Theorem 1.5 (Strong law of large numbers). *Let X_1, X_2, \dots be iid random variables with common mean m . Define $S_n = X_1 + \dots + X_n$. Then $S_n/n \xrightarrow{as} m$.*

The laws of large numbers tell us that the probability mass of an average of random variables “piles up” near its expectation. In just a minute, we will see even more: how fast this piling occurs. But first we should talk about the distribution of the average.

Theorem 1.6 (Central limit theorem). *Let X_1, X_2, \dots be iid with mean μ and variance $\sigma^2 < \infty$. Let $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$. Then,*

$$Z_n := \frac{\bar{X}_n - \mu}{\sqrt{\text{Var}[\bar{X}_n]}} = \frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \rightsquigarrow Z, \quad (1.17)$$

where $Z \sim N(0, 1)$.